

MODEL EXPERIMENTATION

1. SCREENSHOT OF MLFLOW UI BEFORE DROPPING FEATURES

The screenshot shows the MLflow UI interface for experiment 'model_experimentation01'. The left sidebar lists experiments: Default, model_experimentation..., and Lead_scoring. The main panel displays the experiment details, including a description, a table of runs, and various filters. The table shows 11 matching runs with columns for Run Name, Source, Models, and various metrics (AUC, Accuracy, F1, Kappa, MCC, Prec., Recall). The runs are sorted by AUC in descending order.

Run Name	Source	Models	AUC	Accuracy	F1	Kappa	MCC	Prec.	Recall
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702	0.833
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702	0.833
Random For...	ipykernel...	sklearn	0.819	0.737	0.76	0.474	0.482	0.704	0.825
Extra Trees C...	ipykernel...	sklearn	0.818	0.737	0.758	0.473	0.479	0.704	0.82
Decision Tre...	ipykernel...	sklearn	0.817	0.736	0.758	0.472	0.479	0.704	0.82

a. SCREENSHOT OF ALL THE EXPERIMENTS

The screenshot shows the MLflow UI interface for experiment 'model_experimentation01'. The left sidebar lists experiments: Default, model_experimentation..., and Lead_scoring. The main panel displays the experiment details, including a description, a table of runs, and various filters. The table shows 20 matching runs with columns for Run Name, Source, Models, and various metrics (AUC, Accuracy, F1, Kappa, MCC, Prec.). The runs are sorted by AUC in descending order.

Run Name	Source	Models	AUC	Accuracy	F1	Kappa	MCC	Prec.
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702
Light Gradie...	ipykernel...	sklearn	0.821	0.738	0.762	0.476	0.485	0.702
Random For...	ipykernel...	sklearn	0.819	0.737	0.76	0.474	0.482	0.704

b. SCREENSHOT OF ONE EXPERIMENTS WITH ALL THE ARTIFACTS VISIBLE

Light Gradient Boosting Machine

Date: 2023-03-25 22:26:13 Source: `ipykernel_launcher.py` User: root

Status: UNFINISHED Lifecycle Stage: active Parent Run: 9810579c00654fcfb6d399e164d4a9b8

- Description [Edit](#)
- Parameters (20)
- Metrics (8)
- Tags (5)
- Artifacts
 - model
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - AUC.png
 - Confusion Matrix.png
 - Feature Importance.png
 - Holdout.html

Full Path: `/home/mlruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/model` [Register Model](#)

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
------	------

No schema. See [MLflow docs](#) for how to include input and output schema with your model.

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
logged_model = 'runs:/a47a0ba9b11a48f68da4174009656a95/model'

# Load model as a Spark UDF. Override result_type if the model does not return double values.
```

Artifacts

- model
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - AUC.png
 - Confusion Matrix.png
 - Feature Importance.png
 - Holdout.html

Full Path: `/home/mlruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/model` [Register Model](#)

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
------	------

No schema. See [MLflow docs](#) for how to include input and output schema with your model.

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
logged_model = 'runs:/a47a0ba9b11a48f68da4174009656a95/model'

# Load model as a Spark UDF. Override result_type if the model does not return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri=logged_model, result_type='double')

# Predict on a Spark DataFrame.
columns = list(df.columns)
df.withColumn('predictions', loaded_model(*columns)).collect()
```

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/a47a0ba9b11a48f68da4174009656a95/model'

# Load model as a PyFuncModel.
loaded_model = mlflow.pyfunc.load_model(logged_model)

# Predict on a Pandas DataFrame.
```

Light Gradient Boosting Machine

Date: 2023-03-25 22:26:13

Source: ipykernel_launcher.py

User: root

Status: UNFINISHED

Lifecycle Stage: active

Parent Run: [9810579c00654fcb6d399e164d4a9b8](#)

Description [Edit](#)

Parameters (20)

Metrics (8)

Tags (5)

Artifacts

model

MLmodel

conda.yaml

model.pkl

python_env.yaml

requirements.txt

AUC.png

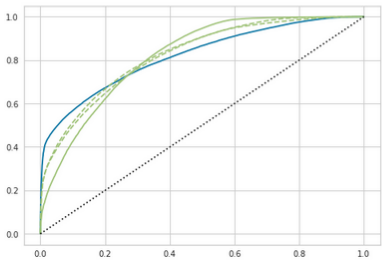
Confusion Matrix.png

Feature Importance.png

Holdout.html

Full Path: /home/miruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/AUC.png

Size: 27.41KB



Light Gradient Boosting Machine

Date: 2023-03-25 22:26:13

Source: ipykernel_launcher.py

User: root

Status: UNFINISHED

Lifecycle Stage: active

Parent Run: [9810579c00654fcb6d399e164d4a9b8](#)

Description [Edit](#)

Parameters (20)

Metrics (8)

Tags (5)

Artifacts

model

MLmodel

conda.yaml

model.pkl

python_env.yaml

requirements.txt

AUC.png

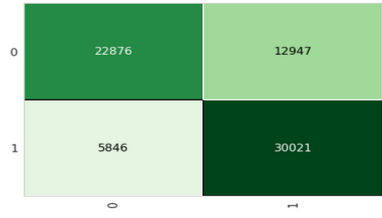
Confusion Matrix.png

Feature Importance.png

Holdout.html

Full Path: /home/miruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/Confusion Matrix.png

Size: 7.6KB



Description [Edit](#)

Parameters (20)

Metrics (8)

Tags (5)

Artifacts

model

MLmodel

conda.yaml

model.pkl

python_env.yaml

requirements.txt

AUC.png

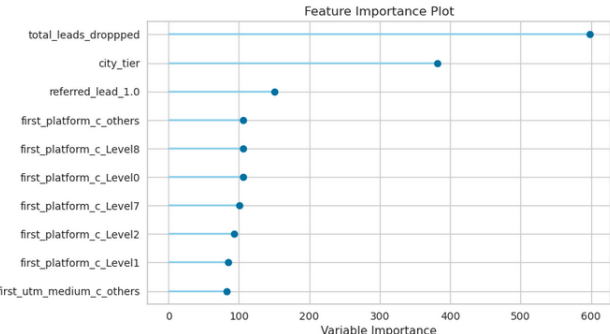
Confusion Matrix.png

Feature Importance.png

Holdout.html

Full Path: /home/miruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/Feature Importance.png

Size: 40.56KB



mlflow

1.26.1

Experiments

Models

model_experimentation01 > Light Gradient Boosting Machine

Light Gradient Boosting Machine

Date: 2023-03-25 22:26:13

Source: ipykernel_launcher.py

User: root

Status: UNFINISHED

Lifecycle Stage: active

Parent Run: 9810579c00654fdb6d399e164d4a9b8

Description [Edit](#)

Parameters (20)

Metrics (8)

Tags (5)

Artifacts

model

MLmodel

conda.yaml

model.pkl

python_env.yaml

requirements.txt

AUC.png

Confusion Matrix.png

Feature Importance.png

Holdout.html

Full Path: /home/miruns/2/a47a0ba9b11a48f68da4174009656a95/artifacts/Holdout.html

Size: 774B

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
0	Light Gradient Boosting Machine	0.7379	0.8197	0.837	0.6987	0.7616	0.4757	0.4853

2. SCREENSHOT OF MLFLOW UI AFTER DROPPING FEATURES

a. SCREENSHOT OF ALL THE EXPERIMENTS

Experiments

model_experimentation02

Share

Search Experiments

Default

model_experimentation...

model_experimentation...

Track machine learning training runs in experiments. [Learn more](#)

Experiment ID: 3

Description [Edit](#)

Refresh

Compare

Delete

Download CSV

↓ AUC

All time

Columns

Only show differences

metrics.rmse < 1 and params.model = "tree"

Search

Filter

Clear

Showing 22 matching runs

	Start Time	Duration	Run Name	User	Source	Version	Models	Metrics >
<input type="checkbox"/>	10 hours ago		Light Gradie...	root	ipykernel...	-	sklearn	↓ AUC 0.821 Accuracy 0.738
<input type="checkbox"/>	10 hours ago		Light Gradie...	root	ipykernel...	-	sklearn	0.821 0.739
<input type="checkbox"/>	10 hours ago		Light Gradie...	root	ipykernel...	-	sklearn	0.821 0.739
<input checked="" type="checkbox"/>	10 hours ago	3.0min	Light Gradie...	root	ipykernel...	-	sklearn	0.821 0.739
<input type="checkbox"/>	10 hours ago		Light Gradie...	root	ipykernel...	-	sklearn	0.821 0.739
<input type="checkbox"/>	10 hours ago		Random For...	root	ipykernel...	-	sklearn	0.819 0.737
<input type="checkbox"/>	10 hours ago		Random For...	root	ipykernel...	-	sklearn	0.819 0.737
<input type="checkbox"/>	10 hours ago		Extra Trees C...	root	ipykernel...	-	sklearn	0.818 0.736

b. SCREENSHOT OF ONE EXPERIMENTS WITH ALL THE ARTIFACTS VISIBLE

Light Gradient Boosting Machine

Date: 2023-03-25 22:49:39 Source: ipykernel_launcher.py User: root

Status: UNFINISHED Lifecycle Stage: active Parent Run: 2280bd2a10284a66b7e96ca71d34e53e

- Description [Edit](#)
- Parameters (20)
- Metrics (8)
- Tags (5)
- Artifacts
 - model
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - AUC.png
 - Confusion Matrix.png
 - Feature Importance.png
 - Holdout.html

Full Path: /home/mlruns/3/01521111f24746a8bfeaa9546aae1224/artifacts/model

[Register Model](#)

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
------	------

No schema. See [MLflow docs](#) for how to include input and output schema with your model.

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
logged_model = 'runs:/01521111f24746a8bfeaa9546aae1224/model'

# Load model as a Spark UDF. Override result_type if the model does not
# return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri=logged_model, re
sult_type='double')

# Predict on a Spark DataFrame.
columns = list(df.columns)
df.withColumn('predictions', loaded_model(*columns)).collect()
```

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/01521111f24746a8bfeaa9546aae1224/model'

# Load model as a PyFuncModel.
```

Light Gradient Boosting Machine

Date: 2023-03-25 22:49:39 Source: ipykernel_launcher.py User: root

Status: UNFINISHED Lifecycle Stage: active Parent Run: 2280bd2a10284a66b7e96ca71d34e53e

- Description [Edit](#)
- Parameters (20)
- Metrics (8)
- Tags (5)
- Artifacts
 - model
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - AUC.png
 - Confusion Matrix.png
 - Feature Importance.png
 - Holdout.html

Full Path: /home/mlruns/3/01521111f24746a8bfeaa9546aae1224/artifacts/model

[Register Model](#)

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
------	------

No schema. See [MLflow docs](#) for how to include input and output schema with your model.

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
logged_model = 'runs:/01521111f24746a8bfeaa9546aae1224/model'

# Load model as a Spark UDF. Override result_type if the model does not
# return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri=logged_model, re
sult_type='double')

# Predict on a Spark DataFrame.
columns = list(df.columns)
df.withColumn('predictions', loaded_model(*columns)).collect()
```

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/01521111f24746a8bfeaa9546aae1224/model'

# Load model as a PyFuncModel.
```

DATA PIPELINE

1. CREATE AN AIRFLOW DAG PYTHON FILE FOR DATA PIPELINE

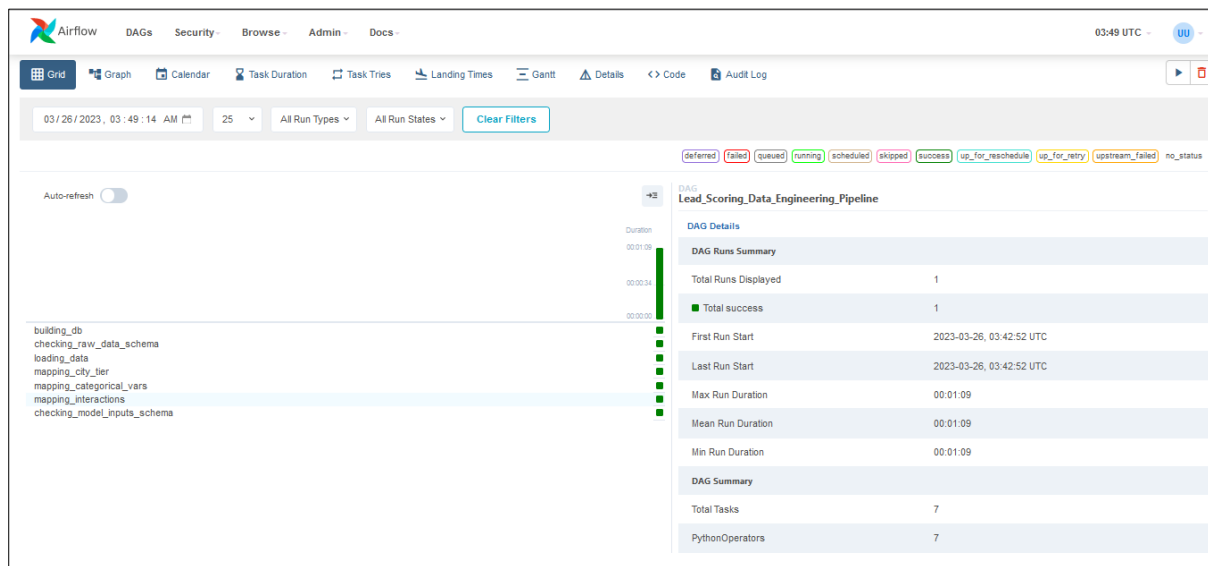
The screenshot shows the Airflow web interface. At the top, there's a navigation bar with links for DAGs, Security, Browse, Admin, and Docs. The current time is 03:42 UTC. Below the navigation bar, there's a section for 'DAG Import Errors (2)' with a warning message about using SQLite as a metadata DB in production. The main section is titled 'DAGs' and shows a list of DAGs. The 'Lead_Scoring_Data_Engineering_Pipeline' DAG is highlighted, showing its owner as 'airflow', its schedule as '@daily', and its last run as '2023-03-24, 00:00:00'.

2. Screenshot of successful execution Airflow DAG in graph

The screenshot shows the Airflow web interface with the 'Lead_Scoring_Data_Engineering_Pipeline' DAG selected. The DAG is in a 'running' state, indicated by a green 'running' badge. The DAG's schedule is '@daily' and its next run is '2023-03-26, 00:00:00'. The DAG's description is 'DAG to run data pipeline for lead scoring'. The DAG's graph is displayed, showing a sequence of tasks: 'building_db', 'checking_raw_data_schema', 'loading_data', 'mapping_city_tier', 'mapping_categorical_vars', 'mapping_interactions', and 'checking_model_inputs_schema'. The 'mapping_interactions' task is currently running, indicated by a green 'running' badge.

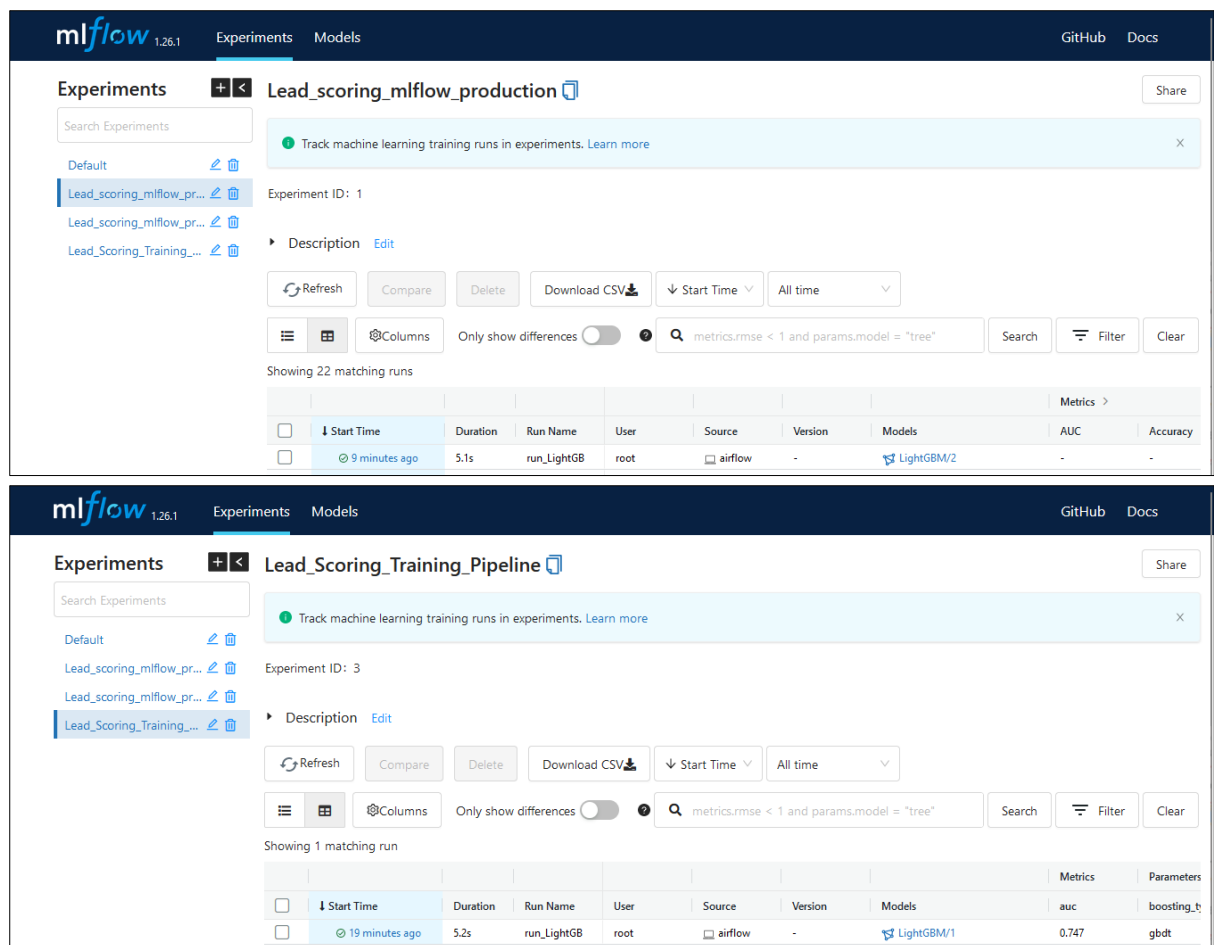
The screenshot shows the Airflow web interface with the 'Lead_Scoring_Data_Engineering_Pipeline' DAG selected. The DAG is in a 'success' state, indicated by a green 'success' badge. The DAG's schedule is '@daily' and its next run is '2023-03-26, 00:00:00'. The DAG's description is 'DAG to run data pipeline for lead scoring'. The DAG's graph is displayed, showing a sequence of tasks: 'building_db', 'checking_raw_data_schema', 'loading_data', 'mapping_city_tier', 'mapping_categorical_vars', 'mapping_interactions', and 'checking_model_inputs_schema'. All tasks are completed, indicated by green 'success' badges.

3. Screenshot of Airflow UI grid



TRAINING PIPELINE

1. Create functions for training pipeline



2. Screenshot of experiments with all the artifacts visible

run_LightGB

Date: 2023-03-26 12:32:24

Duration: 5.1s

Source: airflow

Status: FINISHED

User: root

Lifecycle Stage: active

► Description [Edit](#)

► Parameters (20)

► Metrics (1)

► Tags

▼ Artifacts

models

Full Path:/home/mlruns/1/edcbc89a227544a5a5def75a996f9fa6/artifacts/models

LightGBM, v2
Registered on 2023/03/26

MLmodel
conda.yaml
model.pkl
python_env.yaml
requirements.txt

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. This model is also registered to the [model registry](#).

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
No schema. See MLflow docs for how to include input and output schema with your model.	

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow

# Load model as a Spark UDF. Override result_type if the model does not return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri='logged_model', result_type='double')

# Predict on a Spark DataFrame.
columns = list(df.columns)
df.withColumn('predictions', loaded_model(*columns)).collect()
```

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/edcbc89a227544a5a5def75a996f9fa6/models'

# Load model as a PyFuncModel.
loaded_model = mlflow.pyfunc.load_model(logged_model)

# Predict on a Pandas DataFrame.
import pandas as pd
loaded_model.predict(pd.DataFrame(data))
```

▼ Artifacts

models

Full Path:/home/mlruns/1/edcbc89a227544a5a5def75a996f9fa6/artifacts/models

LightGBM, v2
Registered on 2023/03/26

MLmodel
conda.yaml
model.pkl
python_env.yaml
requirements.txt

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. This model is also registered to the [model registry](#).

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
No schema. See MLflow docs for how to include input and output schema with your model.	

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
logged_model = 'runs:/edcbc89a227544a5a5def75a996f9fa6/models'

# Load model as a Spark UDF. Override result_type if the model does not return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri='logged_model', result_type='double')

# Predict on a Spark DataFrame.
columns = list(df.columns)
df.withColumn('predictions', loaded_model(*columns)).collect()
```

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/edcbc89a227544a5a5def75a996f9fa6/models'

# Load model as a PyFuncModel.
loaded_model = mlflow.pyfunc.load_model(logged_model)

# Predict on a Pandas DataFrame.
import pandas as pd
loaded_model.predict(pd.DataFrame(data))
```

▼ Artifacts

models

Full Path:/home/mlruns/1/edcbc89a227544a5a5def75a996f9fa6/artifacts/models/MLmodel

Size: 434B

```
artifact_path: models
flavors:
  python_function:
    env: conda.yaml
    loader_module: mlflow.sklearn
    model_path: model.pkl
    python_version: 3.8.12
  sklearn:
    code: null
    pickled_model: model.pkl
    serialization_format: cloudpickle
    sklearn_version: 0.23.1
mlflow_version: 1.26.1
model_uri: d72f9a21f2a747f0a9b0a45291a65390e
run_id: edcbc89a227544a5a5def75a996f9fa6
utc_time_created: '2023-03-26 07:02:24.084514'
```


3. Screenshot of model registry with model name and stage as 'production'

The screenshot displays the MLflow Model Registry interface. At the top, there's a navigation bar with 'mlflow 1.26.1', 'Experiments', and 'Models' tabs. Below this, a 'Registered Models' section shows a table of models. The 'LightGBM' model is listed with 'Version 5' in the 'Production' stage. The 'Last Modified' timestamp is '2023-03-26 13:25:25'. Below the table, there are pagination controls showing '1' of 10 pages.

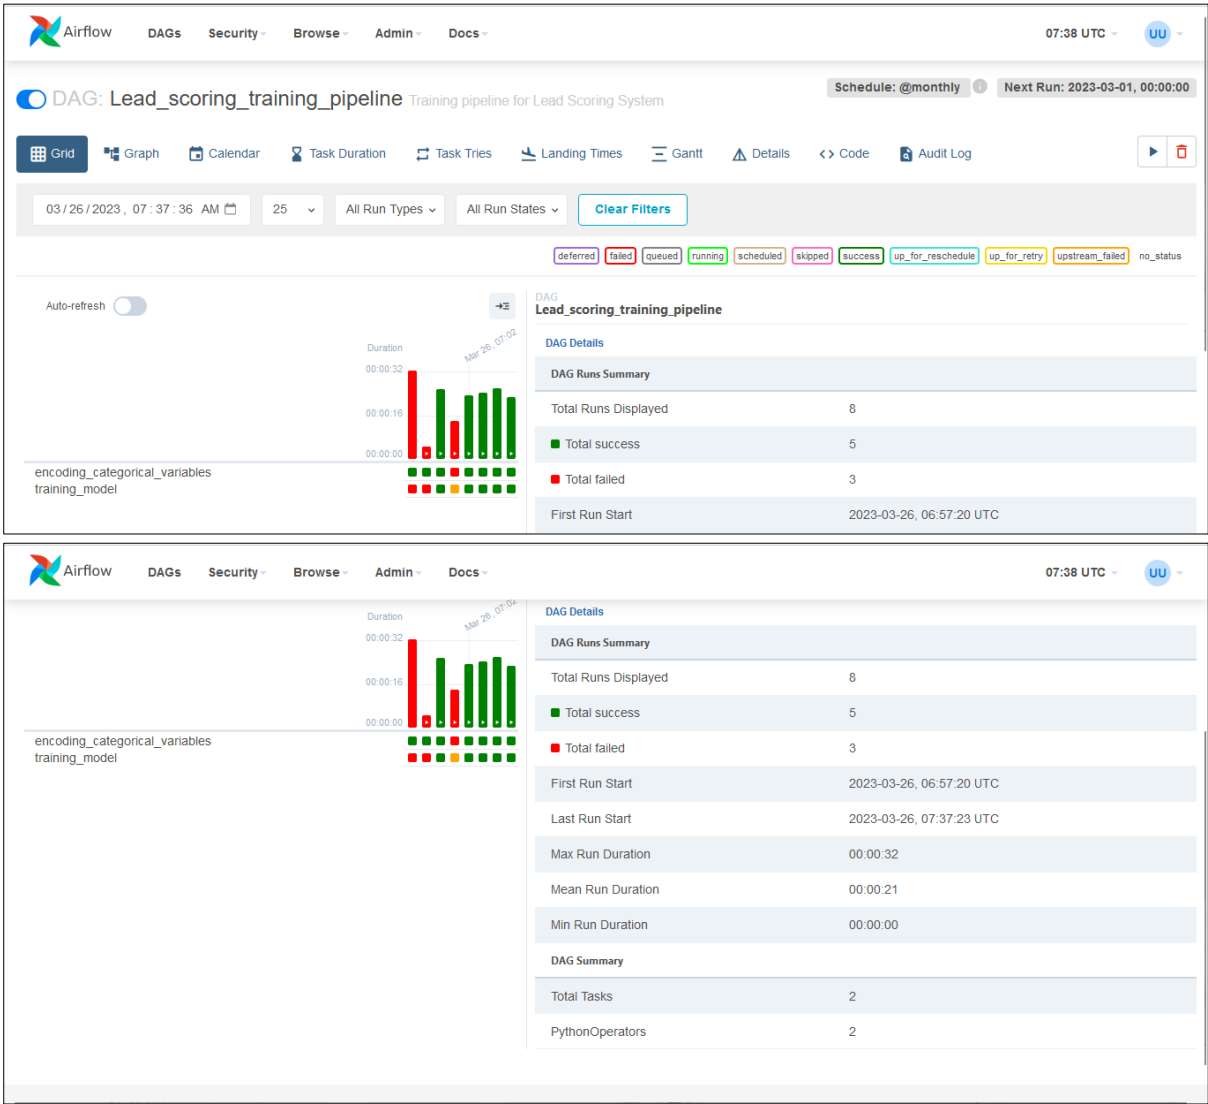
Below the first screenshot, a second screenshot shows the detailed view of the 'LightGBM' model, 'Version 5'. It indicates the model is in the 'Production' stage and was last modified on '2023-03-26 13:25:25'. The 'Source Run' is 'run_LightGB'. The 'Schema' section is empty, with a note: 'No schema. See MLflow docs for how to include input and output schema with your model.'

4. Screenshot of successful execution Airflow DAG in graph

The screenshot shows the Airflow DAG execution interface. At the top, there's a navigation bar with 'Airflow', 'DAGs', 'Security', 'Browse', 'Admin', and 'Docs' tabs. Below this, a 'DAGs' section shows a list of DAGs. The 'Lead_scoring_training_pipeline' is highlighted, showing it is in the 'success' state. The 'Schedule' is '@monthly' and the 'Next Run' is '2023-03-01, 00:00:00'.

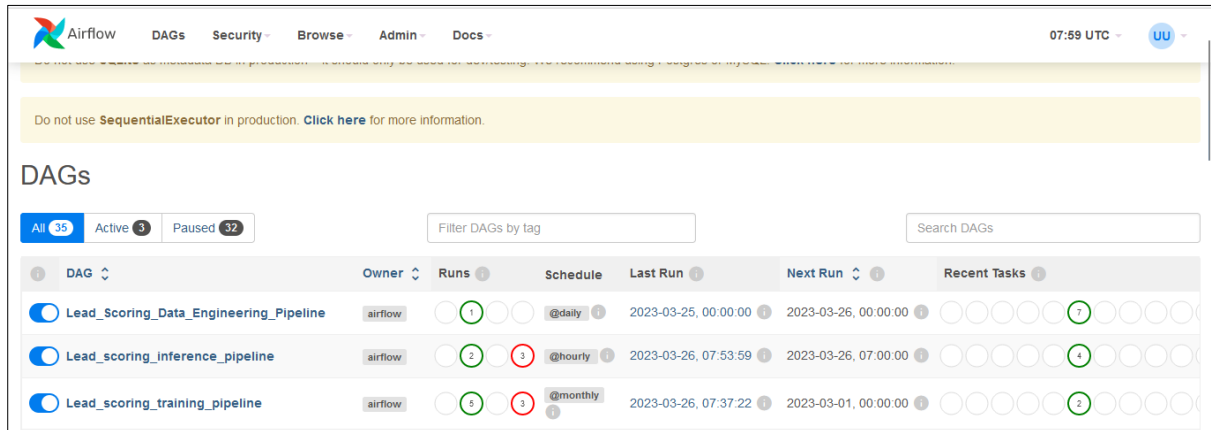
Below the DAG list, a detailed view of the 'Lead_scoring_training_pipeline' is shown. It indicates the DAG is in the 'success' state and the 'Next Run' is '2023-03-01, 00:00:00'. The 'Graph' tab is selected, showing a task dependency graph with two tasks: 'encoding_categorical_variables' and 'training_model'. The 'encoding_categorical_variables' task is in the 'success' state, and the 'training_model' task is in the 'success' state.

5. Screenshot of Airflow UI grid

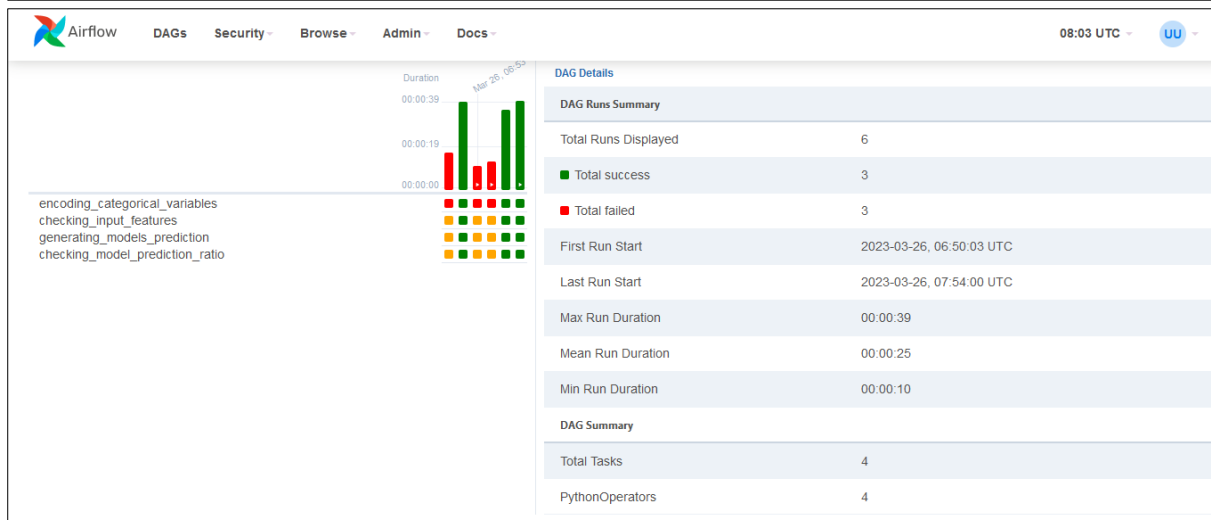
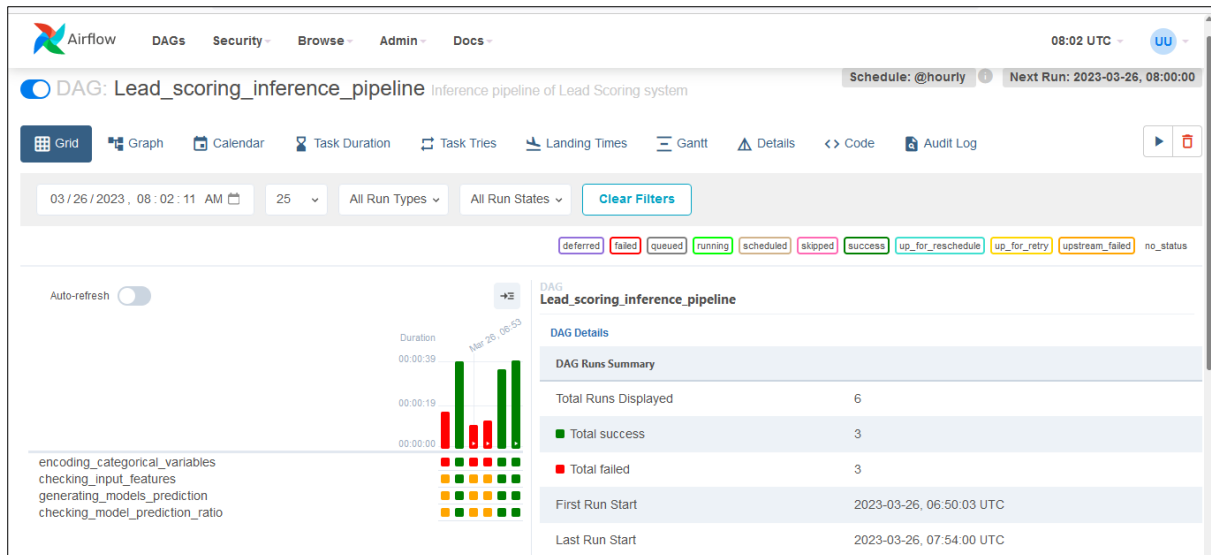


INFERENCE PIPELINE

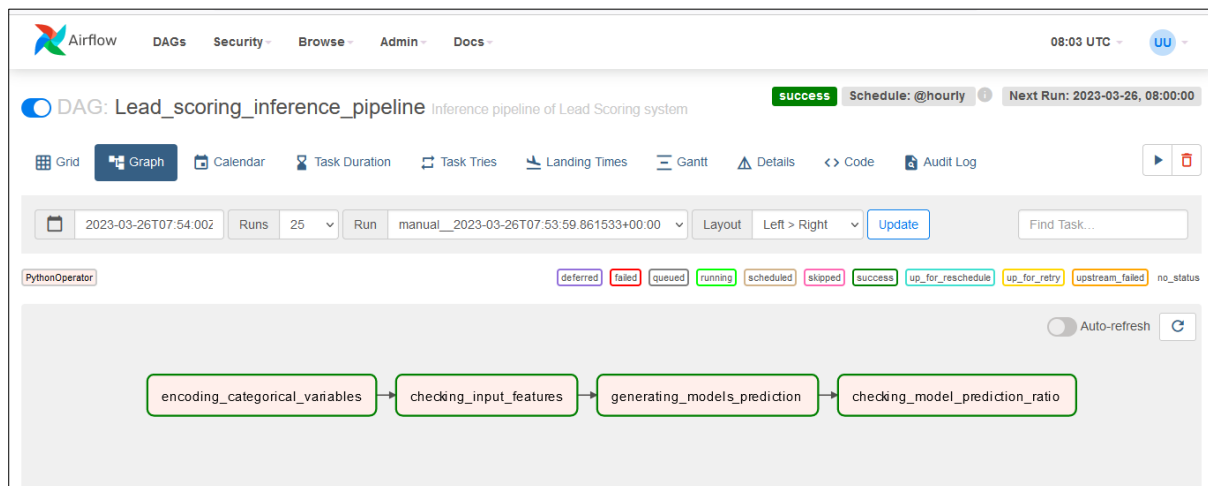
1. Screenshot of Airflow UI DAGS



2. Screenshot of Airflow UI grid



3. Screenshot of successful execution Airflow DAG in graph



Airflow UI for submission

1. Manually trigger the inference pipeline and see if we are getting the expected results

Airflow DAG: Lead_scoring_inference_pipeline Inference pipeline of Lead Scoring system

Grid Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code Audit Log

03/26/2023, 09:05:07 AM 25 All Run Types All Run States Clear Filters

Auto-refresh

encoding_categorical_variables
checking_input_features
generating_models_prediction
checking_model_prediction_ratio

Duration
00:00:43
00:00:21
00:00:00
Mar 26, 07:53

Trigger DAG
Trigger DAG w/ config

DAG Lead_scoring_inference_pipeline

DAG Details

DAG Runs Summary

Total Runs Displayed	9
Total success	6
Total failed	3
First Run Start	2023-03-26, 06:50:03 UTC
Last Run Start	2023-03-26, 09:01:41 UTC
Max Run Duration	00:00:43

Airflow DAG: Lead_scoring_inference_pipeline Inference pipeline of Lead Scoring system

Grid Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code Audit Log

03/26/2023, 09:05:07 AM 25 All Run Types All Run States Clear Filters

Auto-refresh

encoding_categorical_variables
checking_input_features
generating_models_prediction
checking_model_prediction_ratio

Duration
00:00:43
00:00:21
00:00:00
Mar 26, 07:53

DAG Lead_scoring_inference_pipeline

DAG Details

DAG Runs Summary

Total Runs Displayed	9
Total success	6
Total failed	3
First Run Start	2023-03-26, 06:50:03 UTC
Last Run Start	2023-03-26, 09:01:41 UTC

File Edit View Run Kernel Git Tabs Settings Help

Filter files by name

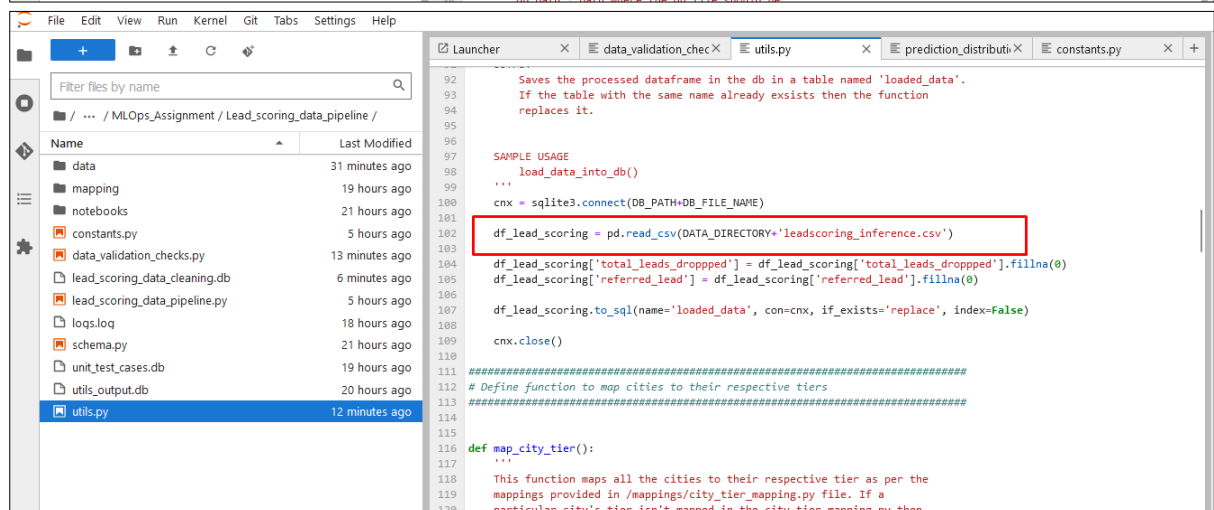
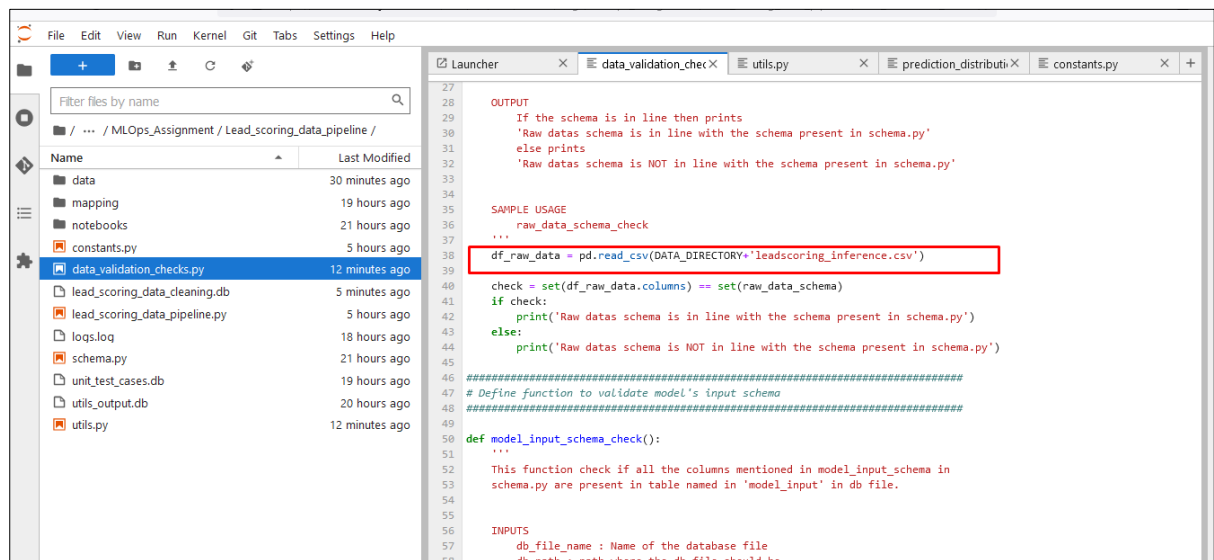
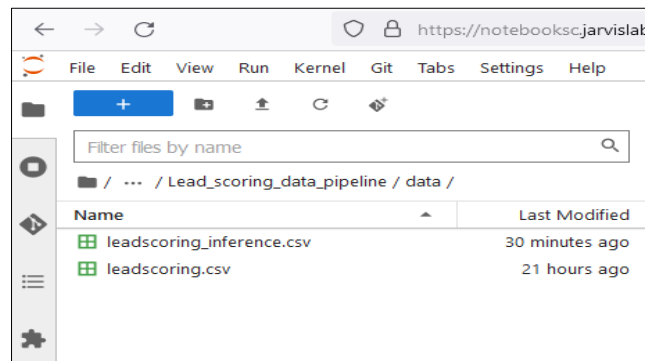
/ ... / ML_Ops_Assignment / Lead_scoring_inference_pipeline /

Name	Last Modified
constants.py	an hour ago
lead_scoring_inference_pipeline.py	an hour ago
prediction_distribution.txt	9 minutes ago
utils.py	an hour ago

Launcher data_validation.py utils.py prediction_distr.py constants.py constants.py

```
1 2023-03-26 07:46:45.437379 %of 1 = 0.5908379504862658 %of 2 = 0.4091620495137343
2 2023-03-26 07:54:38.730560 %of 1 = 0.5908379504862658 %of 2 = 0.4091620495137343
3 2023-03-26 08:00:35.886069 %of 1 = 0.5908379504862658 %of 2 = 0.4091620495137343
4 2023-03-26 08:42:10.867643 %of 1 = 0.5908379504862658 %of 2 = 0.4091620495137343
5 2023-03-26 09:00:22.156069 %of 1 = 0.499719730941704 %of 2 = 0.5002802690582959
6 2023-03-26 09:01:58.257476 %of 1 = 0.499719730941704 %of 2 = 0.5002802690582959
7
```

- Copy the “leadscoring_inference.csv” file present in “Assignment/03_inference_pipeline/Data” to the “airflow/dags/Lead_scoring_data_pipeline/data” folder and modify the utils.py in the Lead_scoring_data_pipeline such that the data pipeline read “leadscoring_inference.csv” instead of “leadscoring.csv”



- Once this is done first trigger the Data pipeline manually and then trigger the inference pipeline once the data pipeline's run is complete. Revise till you get the expected output. Follow the instructions mentioned in the Airflow module to run the Inference pipeline and take the screenshot of Airflow UI for submission. Trigger the training pipeline manually

Do not use **SequentialExecutor** in production. [Click here](#) for more information.

DAGs

All 35 Active 3 Paused 32

Filter DAGs by tag

Search DAGs

DAG	Owner	Runs	Schedule	Last Run	Next Run	Recent Tasks
Lead_Scoring_Data_Engineering_Pipeline	airflow	4	@daily	2023-03-26, 09:12:22	2023-03-26, 00:00:00	7
Lead_scoring_inference_pipeline	airflow	5	@hourly	2023-03-26, 09:01:39	2023-03-26, 09:00:00	4
Lead_scoring_training_pipeline	airflow	5	@monthly	2023-03-26, 08:40:42	2023-03-01, 00:00:00	2

Lead Scoring Data Engineering Pipeline

DAG: Lead_Scoring_Data_Engineering_Pipeline DAG to run data pipeline for lead scoring

Grid Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code Audit Log

03/26/2023, 09:12:24 AM 25 All Run Types All Run States Clear Filters

deferred failed queued running scheduled skipped success up_for_reschedule up_for_retry upstream_failed no_status

Auto-refresh

DAG Lead_Scoring_Data_Engineering_Pipeline

DAG Details

DAG Runs Summary

Run	Start	End	Status
1	2023-03-26, 03:42:52 UTC	2023-03-26, 03:44:00 UTC	success
2	2023-03-26, 04:42:52 UTC	2023-03-26, 04:44:00 UTC	success
3	2023-03-26, 05:42:52 UTC	2023-03-26, 05:44:00 UTC	success
4	2023-03-26, 06:42:52 UTC	2023-03-26, 06:44:00 UTC	running

building_db checking_raw_data_schema loading_data mapping_city_tier mapping_categorical_vars mapping_interactions checking_model_inputs_schema

DAG: Lead_Scoring_Data_Engineering_Pipeline DAG to run data pipeline for lead scoring

Grid Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code Audit Log

2023-03-26T09:12:23Z Runs 25 Run manual_2023-03-26T09:12:22.754578+00:00 Layout Left > Right Update Find Task...

PythonOperator

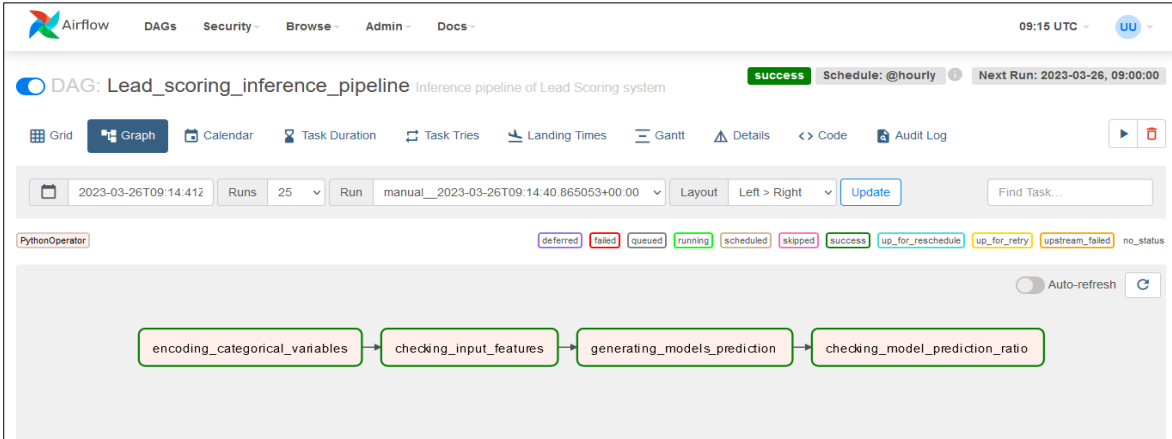
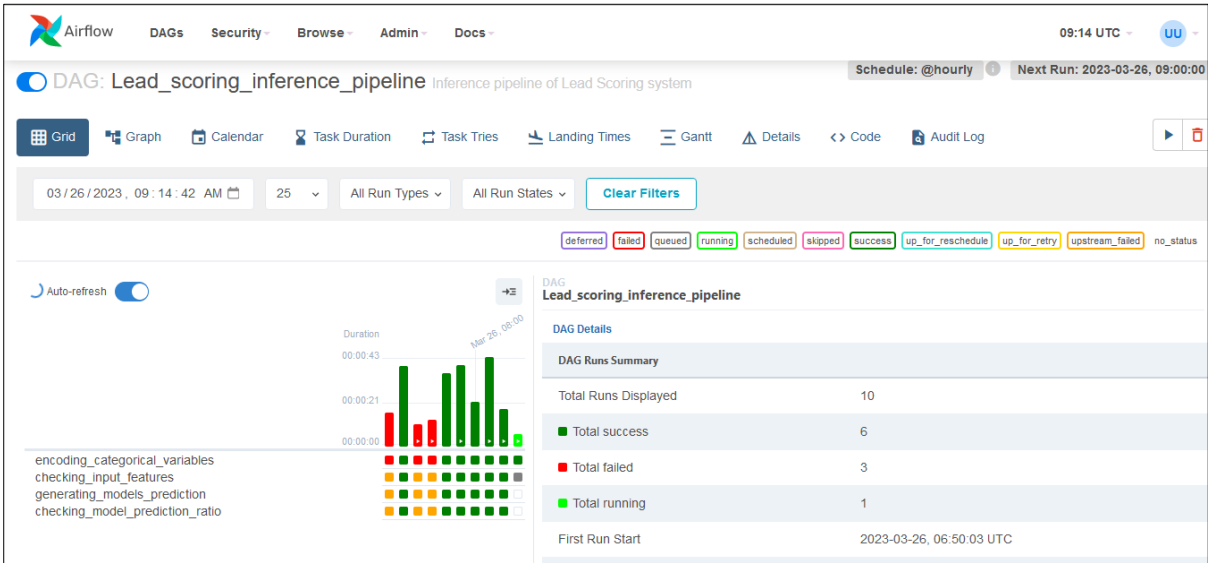
deferred failed queued running scheduled skipped success up_for_reschedule up_for_retry upstream_failed no_status

Auto-refresh

```

graph LR
    building_db --> checking_raw_data_schema
    checking_raw_data_schema --> loading_data
    loading_data --> mapping_city_tier
    mapping_city_tier --> mapping_categorical_vars
    mapping_categorical_vars --> mapping_interactions
    mapping_interactions --> checking_model_inputs_schema
  
```

Lead_scoring_inference_pipeline



Screenshot of Folder structure

Name	Date modified	Type	Size
Lead_scoring_data_pipeline	3/26/2023 3:03 PM	File folder	
Lead_scoring_inference_pipeline	3/26/2023 2:59 PM	File folder	
Lead_scoring_training_pipeline	3/26/2023 3:03 PM	File folder	
notebooks	3/26/2023 2:59 PM	File folder	
unit_test	3/26/2023 3:06 PM	File folder	
Screenshot.pdf	3/26/2023 3:13 PM	Microsoft Edge P...	1,467 KB